

exercise1:

change h_A to d_A in line 54

```
gpuCheck( hipMemcpy(d_A, h_A, bytes, hipMemcpyHostToDevice) );
```

exercise2:

add the code in line 66

```
gpuCheck( hipMemcpy(h_A, d_A, bytes, hipMemcpyDeviceToHost) );
```

exercise3:

add code in line 21

```
if(id<n) A[id]= A[id]* A[id];  
// if use math fuction : pow(A[id],2 ); error as follow  
// Error: h_A[53] = 2808 instead of 2809
```

exercise4:

Kernel is as follow

```
/* -----  
Matrix multiply kernel  
----- */  
__global__ void matrix_multiply(double *A, double *B, double *C, int n)  
{  
    int col = blockDim.x * blockIdx.x + threadIdx.x;  
    int row = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if (col < n && row < n){  
  
        int index = n * row + col;  
        double element = 0.0;  
  
        for (int i=0; i<n; i++){  
  
            int row_index = n * row + i;  
            int col_index = n * i + col;  
  
            element = element + A[row_index] * B[col_index]; //multiply the data  
in Martix A and Martix B  
        }  
  
        C[index]=element; //write the element to C
```

```
}
}
```

exercise5:

通过比较发现使用hipBLAS version of DGEMM，运行的速度更快，效率更高，比用自己实现kenerl完成矩阵乘法的方式提速近8倍。

	Name	Calls	TotalDurationNs	AverageNs	Percentage
1	matrix_multiply(double, double, double*, int)	1	3226886	3226886	88.0429518540558
2	hipBLAS version of DGEMM	1	438241	438241	11.957048145944192

exercise6:

cuda-->hip修改的部分

```
[HIPIFY] info: file './exercises/06_hipify_pingpong/pingpong.cu' statistics:
  CONVERTED refs count: 21
  TOTAL lines of code: 104
  WARNINGS: 0
[HIPIFY] info: CONVERTED refs by names:
  cudaDeviceSynchronize => hipDeviceSynchronize: 1
  cudaError_t => hipError_t: 1
  cudaEventCreate => hipEventCreate: 2
  cudaEventElapsedTime => hipEventElapsedTime: 1
  cudaEventRecord => hipEventRecord: 2
  cudaEventsSynchronize => hipEventsSynchronize: 1
  cudaEvent_t => hipEvent_t: 1
  cudaGetErrorString => hipGetErrorString: 1
  cudaMalloc => hipMalloc: 1
  cudaMallocHost => hipHostMalloc: 1
  cudaMemcpy => hipMemcpy: 4
  cudaMemcpyHostToDevice => hipMemcpyHostToDevice: 4
  cudaSuccess => hipSuccess: 1
```

带宽测试:

H2D最好峰值27.862368899GB/s

D2H最好峰值28.116584723GB/s

```
----- H2D-----
Buffer Size (MiB):    0.007812500, Time (ms):    0.432633996, Bandwidth (GB/s):
0.946758701
Buffer Size (MiB):    0.015625000, Time (ms):    0.462233007, Bandwidth (GB/s):
1.772266341
Buffer Size (MiB):    0.031250000, Time (ms):    0.496313006, Bandwidth (GB/s):
3.301142588
Buffer Size (MiB):    0.062500000, Time (ms):    0.616631985, Bandwidth (GB/s):
5.314028596
Buffer Size (MiB):    0.125000000, Time (ms):    0.773428977, Bandwidth (GB/s):
8.473434793
```

Buffer Size (MiB):	0.250000000	Time (ms):	1.105584979	Bandwidth (GB/s):
11.855443271				
Buffer Size (MiB):	0.500000000	Time (ms):	1.880133986	Bandwidth (GB/s):
13.942836090				
Buffer Size (MiB):	1.000000000	Time (ms):	3.336113930	Bandwidth (GB/s):
15.715530436				
Buffer Size (MiB):	2.000000000	Time (ms):	5.778161049	Bandwidth (GB/s):
18.147226966				
Buffer Size (MiB):	4.000000000	Time (ms):	11.227206230	Bandwidth (GB/s):
18.679197273				
Buffer Size (MiB):	8.000000000	Time (ms):	21.916187286	Bandwidth (GB/s):
19.137927347				
Buffer Size (MiB):	16.000000000	Time (ms):	42.918327332	Bandwidth (GB/s):
19.545514752				
Buffer Size (MiB):	32.000000000	Time (ms):	85.535835266	Bandwidth (GB/s):
19.614254012				
Buffer Size (MiB):	64.000000000	Time (ms):	159.622558594	Bandwidth (GB/s):
21.021108981				
Buffer Size (MiB):	128.000000000	Time (ms):	290.969421387	Bandwidth (GB/s):
23.063888872				
Buffer Size (MiB):	256.000000000	Time (ms):	511.107788086	Bandwidth (GB/s):
26.260160993				
Buffer Size (MiB):	512.000000000	Time (ms):	969.863586426	Bandwidth (GB/s):
27.677650729				
Buffer Size (MiB):	1024.000000000	Time (ms):	1926.867431641	Bandwidth (GB/s):
27.862368899				
----- D2H -----				
Buffer Size (MiB):	0.007812500	Time (ms):	0.445273995	Bandwidth (GB/s):
0.919883048				
Buffer Size (MiB):	0.015625000	Time (ms):	0.466073990	Bandwidth (GB/s):
1.757660839				
Buffer Size (MiB):	0.031250000	Time (ms):	0.515833974	Bandwidth (GB/s):
3.176215765				
Buffer Size (MiB):	0.062500000	Time (ms):	0.586713016	Bandwidth (GB/s):
5.585013304				
Buffer Size (MiB):	0.125000000	Time (ms):	0.786390007	Bandwidth (GB/s):
8.333778336				
Buffer Size (MiB):	0.250000000	Time (ms):	1.171664953	Bandwidth (GB/s):
11.186815791				
Buffer Size (MiB):	0.500000000	Time (ms):	1.964934945	Bandwidth (GB/s):
13.341103259				
Buffer Size (MiB):	1.000000000	Time (ms):	3.307157040	Bandwidth (GB/s):
15.853132879				
Buffer Size (MiB):	2.000000000	Time (ms):	5.848566055	Bandwidth (GB/s):
17.928770746				
Buffer Size (MiB):	4.000000000	Time (ms):	11.176016808	Bandwidth (GB/s):
18.764753455				
Buffer Size (MiB):	8.000000000	Time (ms):	22.144838333	Bandwidth (GB/s):
18.940323415				
Buffer Size (MiB):	16.000000000	Time (ms):	43.046489716	Bandwidth (GB/s):
19.487321859				
Buffer Size (MiB):	32.000000000	Time (ms):	85.300193787	Bandwidth (GB/s):
19.668438318				
Buffer Size (MiB):	64.000000000	Time (ms):	166.399322510	Bandwidth (GB/s):
20.165005178				

Buffer Size (MiB): 128.000000000, Time (ms): 289.394073486, Bandwidth (GB/s): 23.189439642
Buffer Size (MiB): 256.000000000, Time (ms): 508.097656250, Bandwidth (GB/s): 26.415734524
Buffer Size (MiB): 512.000000000, Time (ms): 976.943847656, Bandwidth (GB/s): 27.477060902
Buffer Size (MiB): 1024.000000000, Time (ms): 1909.445678711, Bandwidth (GB/s): 28.116584723

exercise7:

```
/* -----  
Matrix multiply kernel  
----- */  
__global__ void matrix_multiply(double *A, double *B, double *C, int n)  
{  
    __shared__ double s_A[THREADS_PER_BLOCK_Y][THREADS_PER_BLOCK_X];  
    __shared__ double s_B[THREADS_PER_BLOCK_Y][THREADS_PER_BLOCK_X];  
    int col = blockDim.x * blockIdx.x + threadIdx.x;  
    int row = blockDim.y * blockIdx.y + threadIdx.y;  
  
    int lcol = threadIdx.x;  
    int lrow = threadIdx.y;  
  
    int index = n * row + col;  
  
    if (col < n && row < n){  
  
        int THREADS_PER_BLOCK = THREADS_PER_BLOCK_Y;  
        int num_chunks = n / THREADS_PER_BLOCK;  
  
        double element = 0.0;  
  
        for (int chunk=0; chunk<num_chunks; chunk++){  
  
            // TODO: Read data from global GPU memory into shared memory  
            for (int j=0; j<THREADS_PER_BLOCK; j++){  
                s_A[lrow][j]=A[j+chunk*THREADS_PER_BLOCK+col*n];  
                s_B[j][lcol]=B[(j+THREADS_PER_BLOCK*chunk)*n+row];  
            }  
  
            __syncthreads();  
  
            for (int i=0; i<THREADS_PER_BLOCK; i++){  
                element = element + s_A[lrow][i] * s_B[i][lcol];  
            }  
  
            __syncthreads();  
        }  
  
        C[index] = element;  
    }  
}
```

结果:

```
=====
__SUCCESS__
-----
N                : 1024
X Blocks in Grid : 64
X Threads per Block: 16
Y Blocks in Grid : 64
Y Threads per Block: 16
=====
```